

Assignment 3:

Basic Ray Tracing

NAME: CAI YIWEN

STUDENT NUMBER: 2023533145

EMAIL: CAIYW2023@SHANGHAITECH.EDU.CN

ACM Reference Format:

Name: Cai Yiwen, Student Number: 2023533145, Email: caiyw2023@shanghaitech.edu.cn
. 2025. Assignment 3: Basic Ray Tracing. 1, 1 (November 2025), 6 pages.
<https://doi.org/10.1145/nnnnnnnnnnnnnn>

1 INTRODUCTION

This assignment implements a basic ray tracer with geometric acceleration and several illumination models. The mandatory goals are:

- Compile the code base and configure the language server (5%).
- Implement ray-triangle intersection (10%).
- Implement ray-AABB intersection (10%).
- Construct a BVH (Bounding Volume Hierarchy) over primitives (25%).
- Implement IntersectionTestIntegrator and a PerfectRefraction material to validate refractive and solid surface interactions (25%).
- Implement direct lighting with a diffuse BRDF and hard shadow testing (20%).
- Implement anti-aliasing via multi-sampling per pixel (5%).

Optional tasks include rectangular area lights with soft shadows (15%) and environment lighting via environment maps (15%). I completed both optional components by extending the direct lighting integrator and adding an environment-light-based integrator.

2 IMPLEMENTATION DETAILS

2.1 Compilation and Development Environment (must)

I successfully compiled the framework and set up the language server (CMake + clangd) so that all core modules (rdr/accel.h, rdr/integrator.h, rdr/light.h, etc.) are navigable. This allowed fast iteration while editing the BVH, intersection routines, and integrators.

Author's address: Name: Cai Yiwen
Student Number: 2023533145
Email: caiyw2023@shanghaiitech.edu.cn.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM XXXX-XXXX/2025/11-ART
<https://doi.org/10.1145/nnnnnnnnnnnnnn>

2.2 Ray-Triangle Intersection (must)

The function TriangleIntersect implements a Möller–Trumbore style test in accel.cpp. I work in double precision internally:

- Fetch the triangle vertices v_0, v_1, v_2 and ray direction d .
- Compute edges and helper vectors:

$$e_1 = v_1 - v_0, \quad e_2 = v_2 - v_0, \quad s_1 = d \times e_2.$$

- The determinant is $\det = e_1 \cdot s_1$. If $|\det| < \epsilon$, the ray is parallel to the triangle and we return false.
- Otherwise, invert the determinant and compute barycentric coordinates:

$$s = o - v_0, \quad u = (s \cdot s_1) \text{ invDet},$$

$$s_2 = s \times e_1, \quad v = (d \cdot s_2) \text{ invDet},$$

$$t = (e_2 \cdot s_2) \text{ invDet}.$$

- I enforce the barycentric constraints $u \geq 0, v \geq 0, u + v \leq 1$ and range test ray.withinTimeRange(t). Intersections outside $[t_{\min}, t_{\max}]$ are discarded.

If a hit is valid, I call CalculateTriangleDifferentials with barycentric weights $(1 - u - v, u, v)$, assert that interaction.p agrees with ray(t), and shrink ray.t_max to the new, closer intersection.

Listing 1. Ray–triangle intersection (excerpt).

```
InternalVecType e1 = v1 - v0;
InternalVecType e2 = v2 - v0;
InternalVecType ori = Cast<InternalScalarType>(ray.origin);

InternalVecType s1 = Cross(dir, e2);
InternalScalarType det = Dot(e1, s1);
const InternalScalarType eps = 1e-8;
if (abs(det) < eps) return false;

InternalScalarType inv_det = 1 / det;
InternalVecType s = ori - v0;
InternalScalarType u = Dot(s, s1) * inv_det;
if (u < 0 || u > 1) return false;

InternalVecType s2 = Cross(s, e1);
InternalScalarType v = Dot(dir, s2) * inv_det;
if (v < 0 || u + v > 1) return false;

InternalScalarType t = Dot(e2, s2) * inv_det;
if (!ray.withinTimeRange(static_cast<Float>(t))) return
false;
```

2.3 Ray-AABB Intersection (must)

The AABB::intersect method uses the standard “slab” algorithm. For each axis, I compute intersection parameters with the low and high bounds, using the precomputed safe inverse direction:

$$t_0 = (\mathbf{l} - \mathbf{o}) \odot \mathbf{d}^{-1}, \quad t_1 = (\mathbf{u} - \mathbf{o}) \odot \mathbf{d}^{-1}.$$

I then take per-axis min/max:

$$t_{\text{enter}} = \min(t_0, t_1), \quad t_{\text{exit}} = \max(t_0, t_1),$$

and the global entering/exiting times:

$$t_{\text{enter}} = \text{ReduceMax}(t_{\text{enter}}), \quad t_{\text{exit}} = \text{ReduceMin}(t_{\text{exit}}).$$

If $t_{\text{enter}} > t_{\text{exit}}$ or $t_{\text{exit}} < 0$, the ray misses the box. Otherwise, I fill $*t_{\text{in}}$ and $*t_{\text{out}}$ and return true.

Listing 2. Ray-AABB intersection (excerpt).

```
const Vec3f &inv_dir = ray.safe_inverse_direction;
Vec3f origin = ray.origin;
Vec3f t0s = (low_bnd - origin) * inv_dir;
Vec3f t1s = (upper_bnd - origin) * inv_dir;
Vec3f t_enter = Min(t0s, t1s);
Vec3f t_exit = Max(t0s, t1s);
Float t_enter_max = ReduceMax(t_enter);
Float t_exit_min = ReduceMin(t_exit);

if (t_enter_max > t_exit_min || t_exit_min < 0) return false;

*t_in = t_enter_max;
*t_out = t_exit_min;
return true;
```

2.4 BVH Construction (must)

I implemented a generic BVHTree<NodeType> to accelerate intersection:

- Each external NodeType implements getAABB() and getData().
- The internal tree is stored as an array of InternalNode records containing: an AABB, span indices [span_left , span_right], and child indices or leaf flags.
- The build routine takes a span and:
 - (1) Computes the combined AABB over all nodes in the span.
 - (2) Applies a stop criterion: either reaching CUTOFF_DEPTH or having at most two primitives in the span. In that case a leaf node is created.
 - (3) Otherwise, I choose a split dimension as the axis of maximum extent of the AABB via ArgMax(prebuilt_aabb.getExtent()).
 - (4) I compute $\text{split} = \text{span_left} + \text{count}/2$ and use std::nth_element to partition nodes by the centroid coordinate in that dimension:

Listing 3. Median split in BVH build (excerpt).

```
if (depth >= CUTOFF_DEPTH || (span_right - span_left) <= 2)
{
    InternalNode result(span_left, span_right);
    result.is_leaf = true;
    result.aabb = prebuilt_aabb;
    internal_nodes.push_back(result);
```

```
        return internal_nodes.size() - 1;
}

const int &dim = ArgMax(prebuilt_aabb.getExtent());
IndexType count = span_right - span_left;
IndexType split = span_left + count / 2;

auto first = nodes.begin() + span_left;
auto last = nodes.begin() + span_right;
auto nth = nodes.begin() + split;
std::nth_element(first, nth, last,
    [dim](const NodeType &a, const NodeType &b) {
        return a.getAABB().getCenter()[dim] <
            b.getAABB().getCenter()[dim];
});
```

Recursion then builds left and right subtrees and stores the resulting internal node. The intersect routine tests each visited internal node’s AABB using AABB::intersect and recurses only into children whose boxes are hit.

2.5 IntersectionTestIntegrator & PerfectRefraction (must)

The IntersectionTestIntegrator verifies the basic ray tracing logic and materials.

Rendering Loop. The render method iterates over image pixels and uses OpenMP for parallel rows. For each pixel (dx, dy) and each sample:

- (1) Set the sampler’s pixel index: sampler.setPixelIndex2D(Vec2i(dx, dy)).
- (2) Query a random sub-pixel sample position in $[dx, dx + 1) \times [dy, dy + 1)$.
- (3) Generate a differential ray from the camera for that position.
- (4) Call Li(scene, ray, sampler) and accumulate the returned radiance into the film.

Listing 4. IntersectionTestIntegrator::render (excerpt).

```
const Vec2i &resolution = camera->getFilm()->getResolution();

#pragma omp parallel for schedule(dynamic)
for (int dx = 0; dx < resolution.x; dx++) {
    Sampler sampler;
    for (int dy = 0; dy < resolution.y; dy++) {
        sampler.setPixelIndex2D(Vec2i(dx, dy));
        for (int sample = 0; sample < spp; sample++) {
            const Vec2f pixel_sample = sampler.getPixelSample();
            assert(pixel_sample.x >= dx && pixel_sample.x <= dx +
                1);
            assert(pixel_sample.y >= dy && pixel_sample.y <= dy +
                1);
            auto ray = camera->generateDifferentialRay(
                pixel_sample.x, pixel_sample.y);
            const Vec3f &L = Li(scene, ray, sampler);
            camera->getFilm()->commitSample(pixel_sample, L);
        }
    }
}
```

Li Logic. The Li function marches along the ray until it hits a diffuse surface or misses:

- For each bounce, intersect the scene and perform RTTI on `interaction.bsdf` to detect `IdealDiffusion` and `PerfectRefraction`.
- If the material is refractive, I call `bsdf->sample(interaction, sampler, ...)` to get a new incident direction w_i , then spawn a new ray via `interaction.spawnRay(interaction.wi)` and continue.
- If the material is diffuse, I stop the loop and call the direct lighting function.

Listing 5. `IntersectionTestIntegrator::Li` (excerpt).

```
SurfaceInteraction interaction;
bool diffuse_found = false;

for (int i = 0; i < max_depth; ++i) {
    interaction = SurfaceInteraction();
    bool intersected = scene->intersect(ray, interaction);

    bool is_ideal_diffuse =
        dynamic_cast<const IdealDiffusion*>(interaction.bsdf)
        != nullptr;
    bool is_perfect_refraction =
        dynamic_cast<const PerfectRefraction*>(interaction.
        bsdf) != nullptr;

    interaction.wo = -ray.direction;

    if (!intersected) break;

    if (is_perfect_refraction) {
        interaction.bsdf->sample(interaction, sampler, nullptr);
        ray = interaction.spawnRay(interaction.wi);
        continue;
    }

    if (is_ideal_diffuse) {
        diffuse_found = true;
        break;
    }
}

if (!diffuse_found) return Vec3f(0.0f);
return directLighting(scene, interaction);
```

The `PerfectRefraction` material (not shown in full here) computes refraction or total internal reflection based on the incident direction, surface normal, and refractive index ratio, using helper functions like `Refract` and `Reflect`, and stores the new direction in `interaction.wi`.

2.6 Direct Lighting with Diffuse BRDF & Shadows (must)

For the intersection test integrator, I implemented a point-light-based direct lighting routine:

- Compute the vector from the surface point to the point light, its distance, and the normalized light direction.

- Create a “shadow ray” from the surface point toward the light, set `t_max` to slightly less than the distance to the light, and test intersection.
- If any geometry blocks the shadow ray, the point is in shadow and returns black.
- If unoccluded and the BSDF is diffuse, I compute

$$\cos \theta = \max(\mathbf{l} \cdot \mathbf{n}, 0), \quad \text{albedo} = \text{bsdf}-\text{evaluate}(\text{interaction}) \cdot \cos \theta,$$

then apply inverse-square falloff and point-light flux.

Listing 6. Point-light direct lighting (excerpt).

```
Float dist_to_light = Norm(point_light_position -
    interaction.p);
Vec3f light_dir = Normalize(point_light_position -
    interaction.p);
auto test_ray = DifferentialRay(interaction.p, light_dir);
test_ray.setTimeMax(dist_to_light - 1e-3f);

SurfaceInteraction shadow_isect;
if (scene->intersect(test_ray, shadow_isect)) {
    return Vec3f(0.0f);
}

const BSDF *bsdf = interaction.bsdf;
bool is_ideal_diffuse =
    dynamic_cast<const IdealDiffusion*>(bsdf) != nullptr;

if (bsdf && is_ideal_diffuse) {
    Float cos_theta =
        std::max(Dot(light_dir, interaction.normal), 0.0f);
    Vec3f albedo = bsdf->evaluate(interaction) * cos_theta;
    Float inv_r2 = 1.0f / (dist_to_light * dist_to_light + 1e
        -6f);
    return albedo * inv_r2 * point_light_flux;
}
return Vec3f(0.0f);
```

For the BDPT-style direct lighting function, I extended this idea to area lights by sampling points on the `AreaLight` shape, testing occlusion, and weighting contributions by $\cos \theta$, inverse square distance, and the light’s PDF. This produces soft shadows.

2.7 Anti-Aliasing via Multi-Sampling (must)

Anti-aliasing is achieved by shooting multiple rays per pixel. For each pixel, I:

- (1) Initialize the sampler with `setPixelIndex2D`.
- (2) For each of the spp samples, call `getPixelSample()` to obtain a jittered position within the pixel.
- (3) Generate a differential ray and evaluate `Li`.
- (4) Commit the sample back to the film. The film accumulates and later averages the contributions.

This multi-sampling strategy reduces jaggies along edges and produces smoother images.

2.8 Optional: Rectangular Area Lights & Soft Shadows

To support rectangular area lights, I relied on the provided `AreaLight` class and scene description for the Cornell box ceiling light. In `BDPTIntegrator::directLighting`:

- I iterate over all lights in the scene, selecting `AreaLight` instances via `dynamic_cast`.
- For each area light, I draw several samples ($n_{\text{light_samples}} = 8$) by calling `areaLight->sample(interaction, sampler)`, obtaining a light-side `SurfaceInteraction`.
- From the shading point to each sampled light point, I build a shadow ray and test for occlusion. Blocked samples are discarded.
- For visible samples, I compute $\cos \theta$, retrieve emitted radiance L_e , evaluate the diffuse BSDF, apply inverse-square falloff and divide by the area light PDF.

Listing 7. Area-light direct lighting (excerpt).

```
for (const auto &light_ref : lights) {
    const AreaLight *areaLight =
        dynamic_cast<const AreaLight *>(light_ref.get());
    if (!areaLight) continue;

    Vec3f L_light(0.0f);
    const int n_light_samples = 8;
    for (int i = 0; i < n_light_samples; ++i) {
        SurfaceInteraction light_isect =
            areaLight->sample(interaction, sampler);
        Vec3f to_light = light_isect.p - interaction.p;
        Float dist = Norm(to_light);
        Vec3f wi = to_light / dist;
        Float cos_theta = std::max(Dot(wi, interaction.normal),
            0.0f);

        DifferentialRay test_ray(interaction.p, wi);
        test_ray.setTimeMax(dist - 1e-3f);
        SurfaceInteraction shadow_isect;
        if (scene->intersect(test_ray, shadow_isect)) continue;

        Vec3f Le = areaLight->Le(light_isect, -wi);
        const BSDF *bsdf = interaction.bsdf;
        if (!bsdf || !dynamic_cast<const IdealDiffusion *>(bsdf)) continue
            ;

        Float inv_r2 = 1.0f / (dist * dist + 1e-6f);
        L_light += bsdf->evaluate(interaction) * Le *
            cos_theta * inv_r2 / areaLight->pdf(light_isect
                );
    }
    color += L_light / Float(n_light_samples);
}
```

This produces soft penumbras under the Cornell box light, since each shading point sees a different fraction of the emitting rectangle.

2.9 Optional: Environment Lighting via Environment Maps

I implemented an `EnvIntegrator` to handle scenes with environment maps, as configured by an "environment_map" block in JSON. The environment is represented by an `InfiniteAreaLight` retrieved via `scene->getInfiniteLight()`.

`EnvIntegrator::Li`. The logic is:

- (1) Trace the primary ray into the scene.
- (2) If it misses, return the environment radiance in the ray direction: $L = L_{\text{env}}(\mathbf{d})$.
- (3) If it hits a surface, walk through specular refractions similarly to `IntersectionTestIntegrator`, and stop at the first diffuse interaction.
- (4) For a diffuse hit, call `directLighting` that integrates over the environment.

Listing 8. `EnvIntegrator::Li` (excerpt).

```
Vec3f EnvIntegrator::Li(
    ref<Scene> scene, DifferentialRay &ray, Sampler &sampler)
{
    const
    Vec3f color(0.0);
    bool diffuse_found = false;
    SurfaceInteraction interaction;

    for (int depth = 0; depth < max_depth; ++depth) {
        bool intersected = scene->intersect(ray, interaction);
        if (!intersected) {
            const auto &env_light = scene->getInfiniteLight();
            if (env_light) {
                color = env_light->Le(interaction, ray.direction);
            }
            return color;
        }

        bool is_ideal_diffuse =
            dynamic_cast<const IdealDiffusion *>(interaction.bsdf
                ) != nullptr;
        bool is_perfect_refraction =
            dynamic_cast<const PerfectRefraction *>(interaction.
                bsdf) != nullptr;

        interaction.wo = -ray.direction;

        if (is_perfect_refraction) {
            interaction.bsdf->sample(interaction, sampler, nullptr)
                ;
            ray = interaction.spawnRay(interaction.wi);
            continue;
        }

        if (is_ideal_diffuse) {
            diffuse_found = true;
            break;
        }
        break;
    }
}
```

```

if (!diffuse_found) return color;
return directLighting(scene, interaction, sampler);
}

```

Environment Direct Lighting. For environment lighting, I approximate the integral over the upper hemisphere by random sampling:

- Restrict to ideal diffuse surfaces as before.
- Generate random directions by rejection sampling in a cube $[-1, 1]^3$, normalize them, and fold them to lie in the same hemisphere as the surface normal.
- For each direction, cast a shadow ray; if unoccluded, query emitted radiance from the environment light along that direction and accumulate $\text{Le} \times \text{albedo} \times \cos \theta$.
- Average over samples to get the final environment contribution.

Listing 9. Environment direct lighting (excerpt).

```

Vec3f EnvIntegrator::directLighting(
    ref<Scene> scene, SurfaceInteraction &interaction,
    Sampler &sampler) const {
    Vec3f L(0.0f);
    const auto &env_light = scene->getInfiniteLight();
    if (!env_light) return L;

    const BSDF *bsdf = interaction.bsdf;
    bool is_ideal_diffuse =
        dynamic_cast<const IdealDiffusion *>(bsdf) != nullptr;
    if (!bsdf || !is_ideal_diffuse) return L;

    Vec3f albedo = bsdf->evaluate(interaction);
    const int n_samples = 16;
    Vec3f n = interaction.normal;

    for (int i = 0; i < n_samples; ++i) {
        Vec3f dir;
        while (true) {
            Float x = 2.0f * sampler.get1D() - 1.0f;
            Float y = 2.0f * sampler.get1D() - 1.0f;
            Float z = 2.0f * sampler.get1D() - 1.0f;
            dir = Vec3f(x, y, z);
            Float len2 = Dot(dir, dir);
            if (len2 > 1e-4f && len2 <= 1.0f) {
                dir /= std::sqrt(len2);
                break;
            }
        }
        if (Dot(dir, n) < 0.0f) dir = -dir;
        Float cos_theta = Dot(dir, n);
        if (cos_theta <= 0.0f) continue;

        DifferentialRay test_ray(interaction.p, dir);
        SurfaceInteraction shadow_isect;
        if (scene->intersect(test_ray, shadow_isect)) continue;

        SurfaceInteraction dummy;

```


```

        Vec3f Le = env_light->Le(dummy, dir);
        L += Le * albedo * cos_theta;
    }
    L /= Float(n_samples);
    return L;
}


```

On a simple scene with a single diffuse sphere and a constant white environment map, this produces a uniformly lit sphere. When using non-constant environment maps, the visible environment features are correctly reflected in the shading, including occlusions from geometry.

3 RESULTS



(a) Boxes with no refractive material rendering under a single point light
(b) Boxes with refractive material rendering under a single point light



(c) Boxes rendering under a rectangular area light
(d) Sphere rendering under a white environment light

3.1 Qualitative Evaluation

On Cornell box scenes:

3.2 Checklist Against Requirements

- **Must:** Compile the source code and configure the language server (5%) ✓
- **Must:** Ray-triangle intersection (10%) ✓
- **Must:** Ray-AABB intersection (10%) ✓
- **Must:** BVH construction (25%) ✓
- **Must:** IntersectionTestIntegrator and PerfectRefraction material (25%) ✓
- **Must:** Direct lighting with diffuse BRDF and shadow testing (20%) ✓
- **Must:** Anti-aliasing via multi-ray sampling per pixel (5%) ✓
- **Optional:** Rectangular area lights with soft shadows (15%) ✓

6 • Name: Cai Yiwen
Student Number: 2023533145
Email: caiyw2023@shanghaitech.edu.cn

- **Optional:** Environment lighting via environment maps (15%)
✓